

Software Construction

Lab 11

Usman Noor
BESE-5A
Reg# 32273

I Followed following steps to install hadoop
These steps are exclusive for Mac OS

STEP 1: INSTALL HADOOP

```
$ brew search hadoop  
$ brew install hadoop
```

Hadoop will be installed at path /usr/local/Cellar/hadoop

STEP 2: CONFIGURE HADOOP:

Edit hadoop-env.sh, the file can be located at /usr/local/Cellar/hadoop/2.6.0/libexec/etc/hadoop/hadoop-env.sh where 2.6.0 is the hadoop version. Change the line

```
export HADOOP_OPTS="$HADOOP_OPTS -Djava.net.preferIPv4Stack=true"
```

To

```
export HADOOP_OPTS="$HADOOP_OPTS -Djava.net.preferIPv4Stack=true -  
Djava.security.krb5.realm= -Djava.security.krb5.kdc="
```

Edit Core-site.xml, The file can be located at /usr/local/Cellar/hadoop/2.6.0/libexec/etc/hadoop/core-site.xml add below config

```
<property>  
<name>hadoop.tmp.dir</name>  
<value>/usr/local/Cellar/hadoop/hdfs/tmp</value>  
<description>A base for other temporary directories.</description>  
</property>  
<property>  
<name>fs.default.name</name>  
<value>hdfs://localhost:9000</value>  
</property>
```

Edit mapred-site.xml, The file can be located at /usr/local/Cellar/hadoop/2.6.0/libexec/etc/hadoop/mapred-site.xml and by default will be blank add below config

```
<configuration>  
<property>  
<name>mapred.job.tracker</name>  
<value>localhost:9010</value>  
</property>  
</configuration>
```

Edit hdfs-site.xml, The file can be located at /usr/local/Cellar/hadoop/2.6.0/libexec/etc/hadoop/hdfs-site.xml add

```
<configuration>  
<property>  
<name>dfs.replication</name>  
<value></value>
```

```
</property>
</configuration>
```

To simplify life edit a ~/.profile and add the following commands. By default ~/.profile might not exist.

```
alias hstart=<"/usr/local/Cellar/hadoop/2.6.0/sbin/start-dfs.sh;/usr/local/Cellar/hadoop/
2.6.0/sbin/start-yarn.sh">
alias hstop=<"/usr/local/Cellar/hadoop/2.6.0/sbin/stop-yarn.sh;/usr/local/Cellar/hadoop/
2.6.0/sbin/stop-dfs.sh">
```

and source it

```
$ ssh localhost
$ exit
```

STEP 3: RUN HADOOP

```
$ hstart
$ hstop
```

STEP 4: Wordcount: Create a file called WordCount.java.

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;

import java.io.IOException;
import java.util.ArrayList;
import java.util.Iterator;
import java.util.List;
import java.util.StringTokenizer;

public class WordCount extends Configured implements Tool {
    private final static LongWritable ONE = new LongWritable(1L);

    // Mapper Class, Counts words in each line. For each line, break the line into words and
    emits them as (word, 1)

    public static class MapClass extends MapReduceBase implements Mapper<LongWritable,
    Text, Text, IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();
```

```

public void map(LongWritable key, Text value,
    OutputCollector<text, intwritable> output,
    Reporter reporter) throws IOException {

    String line = value.toString();
    StringTokenizer itr = new StringTokenizer(line);
    while (itr.hasMoreTokens()) {
        word.set(itr.nextToken());
        output.collect(word, one);
    }
}

// Reducer class that just emits the sum of the input values.

public static class Reduce extends MapReduceBase implements Reducer< Text,
    IntWritable, Text, IntWritable > {
    public void reduce(Text key, Iterator values,
        OutputCollector<text, intwritable=""> output,
        Reporter reporter) throws IOException {
        int sum = 0;
        while (values.hasNext()) {
            sum += values.next().get();
        }
        output.collect(key, new IntWritable(sum));
    }
}

static int printUsage() {
    System.out.println("wordcount [-m #mappers] [-r #reducers] input_file output_file");
    ToolRunner.printGenericCommandUsage(System.out);
    return -1;
}

public int run(String[] args) throws Exception {

    JobConf conf = new JobConf(getConf(), WordCount.class);
    conf.setJobName("wordcount");

    // the keys are words (strings)
    conf.setOutputKeyClass(Text.class);
    // the values are counts (ints)
    conf.setOutputValueClass(IntWritable.class);

    conf.setMapperClass(MapClass.class);
    // Here we set the combiner!!!!
    conf.setCombinerClass(Reduce.class);
    conf.setReducerClass(Reduce.class);

    List other_args = new ArrayList();
    for(int i=0; i < args.length; ++i) {

```

```

    try {
        if ("-m".equals(args[i])) {
            conf.setNumMapTasks(Integer.parseInt(args[++i]));
        } else if ("-r".equals(args[i])) {
            conf.setNumReduceTasks(Integer.parseInt(args[++i]));
        } else {
            other_args.add(args[i]);
        }
    } catch (NumberFormatException except) {
        System.out.println("ERROR: Integer expected instead of " + args[i]);
        return printUsage();
    } catch (ArrayIndexOutOfBoundsException except) {
        System.out.println("ERROR: Required parameter missing from " +
            args[i-1]);
        return printUsage();
    }
}

// Make sure there are exactly 2 parameters left.
if (other_args.size() != 2) {
    System.out.println("ERROR: Wrong number of parameters: " +
        other_args.size() + " instead of 2.");
    return printUsage();
}

FileInputFormat.setInputPaths(conf, other_args.get(0));
FileOutputFormat.setOutputPath(conf, new Path(other_args.get(1)));

JobClient.runJob(conf);
return 0;
}

public static void main(String[] args) throws Exception {
    int res = ToolRunner.run(new Configuration(), new WordCount(), args);
    System.exit(res);
}
}

```

COMPILE

```
$ javac WordCount.java -cp /usr/local/Cellar/hadoop
```